

## ▼ Bank Customer Churn Prediction

In this project, we use supervised learning models to identify customers who are likely to churn in the future. Furthermore, we will analyze top factors that influence user retention. [Dataset information](#).

## ▼ Contents

- [Part 1: Data Exploration](#)
- [Part 2: Feature Preprocessing](#)
- [Part 3: Model Training and Results Evaluation](#)

## ▼ Part 0: Setup Google Drive Environment / Data Collection

check this [link](#) for more info

```
1 # install pydrive to load data
2 !pip install -U -q PyDrive
3
4 from pydrive.auth import GoogleAuth
5 from pydrive.drive import GoogleDrive
6 from google.colab import auth
7 from oauth2client.client import GoogleCredentials
8
9 auth.authenticate_user()
10 gauth = GoogleAuth()
11 gauth.credentials = GoogleCredentials.get_application_default()
12 drive = GoogleDrive(gauth)
```

```
1 # the same way we get id from last class
2 #https://drive.google.com/file/d/1szdCZ98EK59cfJ4jG03g1H0v_0hC1oyN/view?usp=sharing
3 id = "1szdCZ98EK59cfJ4jG03g1H0v_0hC1oyN"
4 file = drive.CreateFile({'id':id})
5 file.GetContentFile('bank_churn.csv')
```

```
1 import pandas as pd
2
3 df = pd.read_csv('bank_churn.csv')
4 df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.0
1	2	15647311	Hill	608	Spain	Female	41	1	83807.1
2	3	15619304	Onio	502	France	Female	42	8	159660.1
3	4	15701354	Boni	699	France	Female	39	1	0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.1

## ▼ Part 1: Data Exploration

### ▼ Part 1.1: Understand the Raw Dataset

```
1 import pandas as pd
2 import numpy as np
3
4 churn_df = pd.read_csv('bank_churn.csv')
```

```
1 churn_df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.0
1	2	15647311	Hill	608	Spain	Female	41	1	83807.1
2	3	15619304	Onio	502	France	Female	42	8	159660.1
3	4	15701354	Boni	699	France	Female	39	1	0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.1

```
1 # check data info
2 churn_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RowNumber       10000 non-null  int64
1   CustomerId      10000 non-null  int64
2   Surname         10000 non-null  object
3   CreditScore     10000 non-null  int64
```

```

4 Geography      10000 non-null object
5 Gender         10000 non-null object
6 Age           10000 non-null int64
7 Tenure        10000 non-null int64
8 Balance       10000 non-null float64
9 NumOfProducts 10000 non-null int64
10 HasCrCard     10000 non-null int64
11 IsActiveMember 10000 non-null int64
12 EstimatedSalary 10000 non-null float64
13 Exited        10000 non-null int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

```

1 # check the unique values for each column
2 churn_df.nunique()

```

```

RowNumber      10000
CustomerId     10000
Surname         2932
CreditScore     460
Geography        3
Gender           2
Age             70
Tenure          11
Balance        6382
NumOfProducts    4
HasCrCard        2
IsActiveMember   2
EstimatedSalary 9999
Exited          2
dtype: int64

```

```

1 # Get target variable
2 y = churn_df['Exited']

```

## ▼ Part 1.2: Understand the features

```

1 # check missing values
2 churn_df.isnull().sum()

```

```

RowNumber      0
CustomerId     0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0

```

```
IsActiveMember    0
EstimatedSalary   0
Exited            0
dtype: int64
```

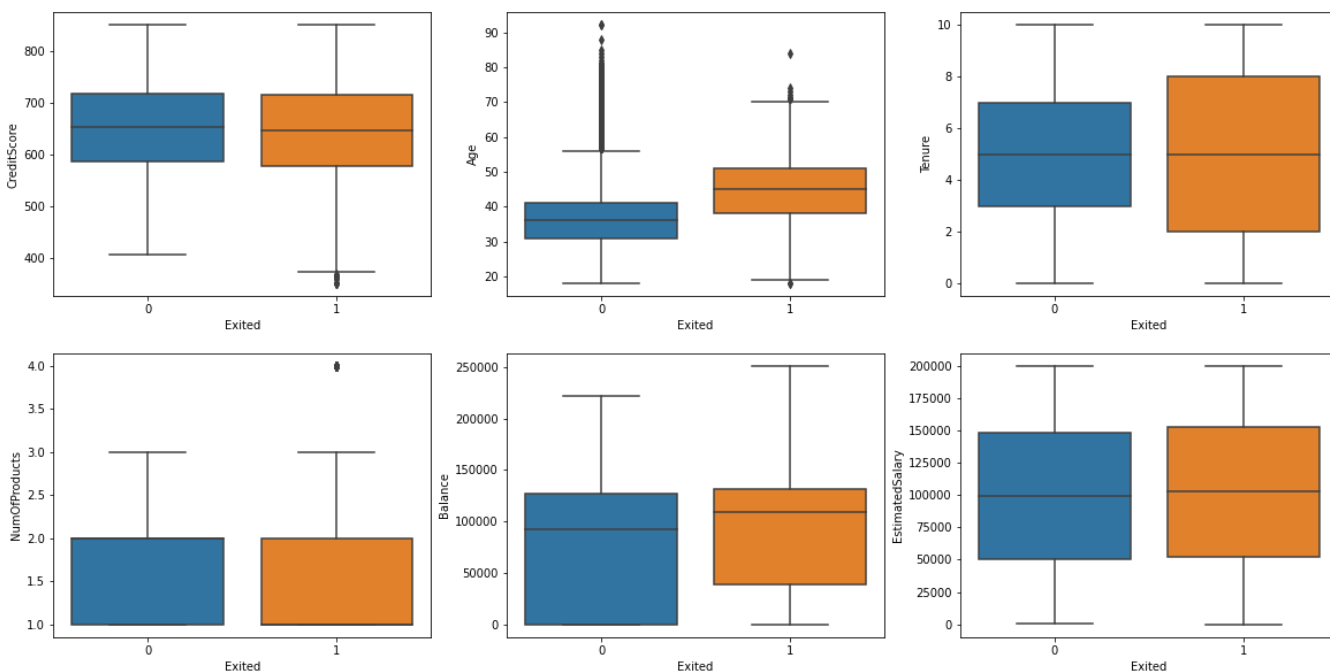
```
1 # understand Numerical feature
2 # discrete/continuous
3 # 'CreditScore', 'Age', 'Tenure', 'NumberOfProducts'
4 # 'Balance', 'EstimatedSalary'
5 churn_df[['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'EstimatedSalary']].d
```

	CreditScore	Age	Tenure	NumOfProducts	Balance	EstimatedSalary
<b>count</b>	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
<b>mean</b>	650.528800	38.921800	5.012800	1.530200	76485.889288	100090.200000
<b>std</b>	96.653299	10.487806	2.892174	0.581654	62397.405202	57510.400000
<b>min</b>	350.000000	18.000000	0.000000	1.000000	0.000000	11.500000
<b>25%</b>	584.000000	32.000000	3.000000	1.000000	0.000000	51002.100000
<b>50%</b>	652.000000	37.000000	5.000000	1.000000	97198.540000	100193.900000
<b>75%</b>	718.000000	44.000000	7.000000	2.000000	127644.240000	149388.200000
<b>max</b>	850.000000	92.000000	10.000000	4.000000	250898.090000	199992.400000

```
1 # check the feature distribution
2 # pandas.DataFrame.describe()
3 # boxplot, distplot, countplot
4 import matplotlib.pyplot as plt
5 import seaborn as sns
```

```
1 # boxplot for numerical feature
2 _,axss = plt.subplots(2,3, figsize=[20,10])
3 sns.boxplot(x='Exited', y='CreditScore', data=churn_df, ax=axss[0][0])
4 sns.boxplot(x='Exited', y='Age', data=churn_df, ax=axss[0][1])
5 sns.boxplot(x='Exited', y='Tenure', data=churn_df, ax=axss[0][2])
6 sns.boxplot(x='Exited', y='NumOfProducts', data=churn_df, ax=axss[1][0])
7 sns.boxplot(x='Exited', y='Balance', data=churn_df, ax=axss[1][1])
8 sns.boxplot(x='Exited', y='EstimatedSalary', data=churn_df, ax=axss[1][2])
```

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fb3ab028c10&gt;




```

1 # correlations between features
2 corr_score = churn_df[['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'EstimatedSalary']]
3
4 # show heatmap of correlations
5 sns.heatmap(corr_score)

```

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fb3aae68ad0&gt;





```
1 # check the actual values of correlations
2 corr_score
```

	CreditScore	Age	Tenure	NumOfProducts	Balance	EstimatedSalary
CreditScore	1.000000	-0.003965	0.000842	0.012238	0.006268	-0.001384
Age	-0.003965	1.000000	-0.009997	-0.030680	0.028308	-0.007201
Tenure	0.000842	-0.009997	1.000000	0.013444	-0.012254	0.007784
NumOfProducts	0.012238	-0.030680	0.013444	1.000000	-0.304180	0.014204
Balance	0.006268	0.028308	-0.012254	-0.304180	1.000000	0.012797
EstimatedSalary	-0.001384	-0.007201	0.007784	0.014204	0.012797	1.000000

```
1 # understand categorical feature
2 # 'Geography', 'Gender'
3 # 'HasCrCard', 'IsActiveMember'
4 _,axss = plt.subplots(2,2, figsize=[20,10])
5 sns.countplot(x='Exited', hue='Geography', data=churn_df, ax=axss[0][0])
6 sns.countplot(x='Exited', hue='Gender', data=churn_df, ax=axss[0][1])
7 sns.countplot(x='Exited', hue='HasCrCard', data=churn_df, ax=axss[1][0])
8 sns.countplot(x='Exited', hue='IsActiveMember', data=churn_df, ax=axss[1][1])
```

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fb3aac7fcd0&gt;



## Part 2: Feature Preprocessing

```
1 # Get feature space by dropping useless feature
2 to_drop = ['RowNumber', 'CustomerId', 'Surname', 'Exited']
3 X = churn_df.drop(to_drop, axis=1)
```

```
1 X.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

```
1 X.dtypes
```

```
CreditScore      int64
Geography         object
Gender            object
Age              int64
Tenure           int64
Balance          float64
NumOfProducts    int64
HasCrCard        int64
IsActiveMember   int64
EstimatedSalary  float64
dtype: object
```

```
1 cat_cols = X.columns[X.dtypes == 'O']
2 num_cols = X.columns[(X.dtypes == 'float64') | (X.dtypes == 'int64')]
```

```
1 num_cols
```

```
Index(['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary'],
      dtype='object', length=8)
```

```
dtype='object')
```

```
1 cat_cols
```

```
Index(['Geography', 'Gender'], dtype='object')
```

## Split dataset

```
1 # Split data into training and testing
2 # 100 -> 75:y=1, 25:y=0
3 # training(80): 60 y=1; 20 y=0
4 # testing(20): 15 y=1; 5 y=0
5
6 from sklearn import model_selection
7
8 # Reserve 25% for testing
9 # stratify example:
10 # 100 -> y: 80 '0', 20 '1' -> 4:1
11 # 80% training 64: '0', 16:'1' -> 4:1
12 # 20% testing 16:'0', 4: '1' -> 4:1
13 X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.25,
14
15 print('training data has ' + str(X_train.shape[0]) + ' observation with ' + str(X_train.sh
16 print('test data has ' + str(X_test.shape[0]) + ' observation with ' + str(X_test.shape[1]
```

```
training data has 7500 observation with 10 features
test data has 2500 observation with 10 features
```

- 10000 -> 8000 '0' + 2000 '1'
- 25% test 75% training

---

without stratified sampling:

### ▼ • extreme case:

1. testing: 2000 '1' + 500 '0'
2. training: 7500 '0'

---

with stratified sampling:

1. testing: 2000 '0' + 500 '1'
2. training: 6000 '0' + 1500 '1'

Read more for handling [categorical feature](#), and there is an awesome package for [encoding](#).



```
1 X_train.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	
<b>7971</b>	633	Spain	Male	42	10	0.00	1	0	
<b>9152</b>	708	Germany	Female	23	4	71433.08	1	1	
<b>6732</b>	548	France	Female	37	9	0.00	2	0	
<b>902</b>	645	France	Female	48	7	90612.34	1	1	
<b>2996</b>	729	Spain	Female	45	7	91091.06	2	1	

```

1 # One hot encoding
2 # another way: get_dummies
3 from sklearn.preprocessing import OneHotEncoder
4
5 def OneHotEncoding(df, enc, categories):
6     transformed = pd.DataFrame(enc.transform(df[categories]).toarray(), columns=enc.get_feature_names_out(categories))
7     return pd.concat([df.reset_index(drop=True), transformed], axis=1).drop(categories, axis=1)
8
9 categories = ['Geography']
10 enc_ohe = OneHotEncoder()
11 enc_ohe.fit(X_train[categories])
12
13 X_train = OneHotEncoding(X_train, enc_ohe, categories)
14 X_test = OneHotEncoding(X_test, enc_ohe, categories)
15

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: The class OneHotEncoder is deprecated in version 0.22 and will be removed in version 0.24. Use OneHotEncoder from sklearn.preprocessing instead.
warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: The class OneHotEncoder is deprecated in version 0.22 and will be removed in version 0.24. Use OneHotEncoder from sklearn.preprocessing instead.
warnings.warn(msg, category=FutureWarning)

```

```
1 X_train.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
<b>0</b>	633	Male	42	10	0.00	1	0	1
<b>1</b>	708	Female	23	4	71433.08	1	1	0
<b>2</b>	548	Female	37	9	0.00	2	0	0
<b>3</b>	645	Female	48	7	90612.34	1	1	1
<b>4</b>	729	Female	45	7	91091.06	2	1	0

```

1 # Ordinal encoding
2 from sklearn.preprocessing import OrdinalEncoder
3
4 categories = ['Gender']
5 enc_oe = OrdinalEncoder()
6 enc_oe.fit(X_train[categories])
7
8 X_train[categories] = enc_oe.transform(X_train[categories])
9 X_test[categories] = enc_oe.transform(X_test[categories])

```

```
1 X_train.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	633	1.0	42	10	0.00	1	0	1
1	708	0.0	23	4	71433.08	1	1	0
2	548	0.0	37	9	0.00	2	0	0
3	645	0.0	48	7	90612.34	1	1	1
4	729	0.0	45	7	91091.06	2	1	0

## Standardize/Normalize Data

```

1 # Scale the data, using standardization
2 # standardization (x-mean)/std
3 # normalization (x-x_min)/(x_max-x_min) ->[0,1]
4
5 # 1. speed up gradient descent
6 # 2. same scale
7 # 3. algorithm requirments
8
9 # for example, use training data to train the standardscaler to get mean and std
10 # apply mean and std to both training and testing data.
11 # fit_transform does the training and applying, transform only does applying.
12 # Because we can't use any info from test, and we need to do the same modification
13 # to testing data as well as training data
14
15 # https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_scaling.html#sphx-g
16 # https://scikit-learn.org/stable/modules/preprocessing.html
17
18
19 # min-max example: (x-x_min)/(x_max-x_min)
20 # [1,2,3,4,5,6,100] -> fit(min:1, max:6) (scalar.min = 1, scalar.max = 6) -> transform [(1
21 # scalar.fit(train) -> min:1, max:100
22 # scalar.transform(apply to x) -> apply min:1, max:100 to X_train
23 # scalar.transform -> apply min:1, max:100 to X_test
24

```

```

25 # scalar.fit -> mean:1, std:100
26 # scalar.transform -> apply mean:1, std:100 to X_train
27 # scalar.transform -> apply mean:1, std:100 to X_test
28
29 from sklearn.preprocessing import StandardScaler
30 scaler = StandardScaler()
31 scaler.fit(X_train[num_cols])
32 X_train[num_cols] = scaler.transform(X_train[num_cols])
33 X_test[num_cols] = scaler.transform(X_test[num_cols])

1 X_train.head()

```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	633	1.0	42	10	0.00	1	0	1
1	708	0.0	23	4	71433.08	1	1	0
2	548	0.0	37	9	0.00	2	0	0
3	645	0.0	48	7	90612.34	1	1	1
4	729	0.0	45	7	91091.06	2	1	0

## ▼ Part 3: Model Training and Result Evaluation

### ▼ Part 3.1: Model Training

```

1 #@title build models
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.linear_model import LogisticRegression
5
6 # Logistic Regression
7 classifier_logistic = LogisticRegression()
8
9 # K Nearest Neighbors
10 classifier_KNN = KNeighborsClassifier()
11
12 # Random Forest
13 classifier_RF = RandomForestClassifier()

```

build models



```

1 # Train the model
2 classifier_logistic.fit(X_train, y_train)

```

LogisticRegression()

```
1 # Prediction of test data
2 classifier_logistic.predict(X_test)
```

```
array([0, 0, 0, ..., 0, 0, 0])
```

```
1 # Accuracy of test data
2 classifier_logistic.score(X_test, y_test)
```

```
0.7936
```

```
1 # Use 5-fold Cross Validation to get the accuracy for different models
2 model_names = ['Logistic Regression', 'KNN', 'Random Forest']
3 model_list = [classifier_logistic, classifier_KNN, classifier_RF]
4 count = 0
5
6 for classifier in model_list:
7     cv_score = model_selection.cross_val_score(classifier, X_train, y_train, cv=5)
8     print(cv_score)
9     print('Model accuracy of ' + model_names[count] + ' is ' + str(cv_score.mean()))
10    count += 1
```

```
[0.78533333 0.78933333 0.78866667 0.792      0.78933333]
Model accuracy of Logistic Regression is 0.7889333333333333
[0.76333333 0.772      0.75733333 0.76466667 0.758      ]
Model accuracy of KNN is 0.7630666666666668
[0.88133333 0.864      0.85266667 0.85533333 0.86533333]
Model accuracy of Random Forest is 0.8637333333333332
```

## ▼ (Optional) Part 3.2: Use Grid Search to Find Optimal Hyperparameters

alternative: random search

```
1 #Loss/cost function --> (wx + b - y) ^2 + λ * |w| --> λ is a hyperparameter
```

```
1 from sklearn.model_selection import GridSearchCV
2
3 # helper function for printing out grid search results
4 def print_grid_search_metrics(gs):
5     print ("Best score: " + str(gs.best_score_))
6     print ("Best parameters set:")
7     best_parameters = gs.best_params_
8     for param_name in sorted(best_parameters.keys()):
9         print(param_name + ':' + str(best_parameters[param_name]))
```

### ▼ Part 3.2.1: Find Optimal Hyperparameters - LogisticRegression

```

1 # Possible hyperparamter options for Logistic Regression Regularization
2 # Penalty is choosed from L1 or L2
3 # C is the 1/lambda value(weight) for L1 and L2
4 # solver: algorithm to find the weights that minimize the cost function
5
6 # ('l1', 0.01)('l1', 0.05) ('l1', 0.1) ('l1', 0.2)('l1', 1)
7 # ('l2', 0.01)('l2', 0.05) ('l2', 0.1) ('l2', 0.2)('l2', 1)
8 parameters = {
9     'penalty':('l1', 'l2'),
10    'C':(0.01, 0.05, 0.1, 0.2, 1)
11 }
12 Grid_LR = GridSearchCV(LogisticRegression(solver='liblinear'),parameters, cv=5)
13 Grid_LR.fit(X_train, y_train)

```

```

GridSearchCV(cv=5, estimator=LogisticRegression(solver='liblinear'),
             param_grid={'C': (0.01, 0.05, 0.1, 0.2, 1),
                          'penalty': ('l1', 'l2')})

```

```

1 # the best hyperparameter combination
2 # C = 1/lambda
3 print_grid_search_metrics(Grid_LR)

```

```

Best score: 0.8130666666666666
Best parameters set:
C:0.05
penalty:l1

```

```

1 # best model
2 best_LR_model = Grid_LR.best_estimator_

```

```

1 best_LR_model.predict(X_test)

array([0, 0, 0, ..., 0, 0, 0])

```

```

1 best_LR_model.score(X_test, y_test)

0.8084

```

```

1 LR_models = pd.DataFrame(Grid_LR.cv_results_)
2 res = (LR_models.pivot(index='param_penalty', columns='param_C', values='mean_test_score')
3        )
4 _ = sns.heatmap(res, cmap='viridis')

```



### ▼ Part 3.2.2: Find Optimal Hyperparameters: KNN

```
1 # Possible hyperparamter options for KNN
2 # Choose k
3 parameters = {
4     'n_neighbors':[1,3,5,7,9]
5 }
6 Grid_KNN = GridSearchCV(KNeighborsClassifier(),parameters, cv=5)
7 Grid_KNN.fit(X_train, y_train)
```

```
GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
              param_grid={'n_neighbors': [1, 3, 5, 7, 9]})
```

```
1 # best k
2 print_grid_search_metrics(Grid_KNN)
```

```
Best score: 0.7856
Best parameters set:
n_neighbors:9
```

```
1 best_KNN_model = Grid_KNN.best_estimator_
```

### ▼ Part 3.2.3: Find Optimal Hyperparameters: Random Forest

```
1 # Possible hyperparamter options for Random Forest
2 # Choose the number of trees
3 parameters = {
4     'n_estimators' : [60,80,100],
5     'max_depth': [1,5,10]
6 }
7 Grid_RF = GridSearchCV(RandomForestClassifier(),parameters, cv=5)
8 Grid_RF.fit(X_train, y_train)
```

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(),
              param_grid={'max_depth': [1, 5, 10],
                           'n_estimators': [60, 80, 100]})
```

```
1 # best number of tress
2 print_grid_search_metrics(Grid_RF)
```

```
Best score: 0.8667999999999999
Best parameters set:
max_depth:10
n_estimators:100
```

```
1 # best random forest
2 best_RF_model = Grid_RF.best_estimator_
```

```
1 best_RF_model
```

```
RandomForestClassifier(max_depth=10)
```

### ▼ Part 3.3: Model Evaluation - Confusion Matrix (Precision, Recall, Accuracy)

class of interest as positive

TP: correctly labeled real churn

Precision (PPV, positive predictive value):  $tp / (tp + fp)$ ; Total number of true predictive churn divided by the total number of predictive churn; High Precision means low fp, not many return users were predicted as churn users.

Recall (sensitivity, hit rate, true positive rate):  $tp / (tp + fn)$  Predict most positive or churn user correctly. High recall means low fn, not many churn users were predicted as return users.

```
1 from sklearn.metrics import confusion_matrix
2 from sklearn.metrics import classification_report
3 from sklearn.metrics import precision_score
4 from sklearn.metrics import recall_score
5
6 # calculate accuracy, precision and recall, [[tn, fp],[]]
7 def cal_evaluation(classifier, cm):
8     tn = cm[0][0]
9     fp = cm[0][1]
10    fn = cm[1][0]
11    tp = cm[1][1]
12    accuracy = (tp + tn) / (tp + fp + fn + tn + 0.0)
13    precision = tp / (tp + fp + 0.0)
14    recall = tp / (tp + fn + 0.0)
15    print (classifier)
16    print ("Accuracy is: " + str(accuracy))
17    print ("precision is: " + str(precision))
18    print ("recall is: " + str(recall))
19    print ()
20
```

```

21 # print out confusion matrices
22 def draw_confusion_matrices(confusion_matrices):
23     class_names = ['Not', 'Churn']
24     for cm in confusion_matrices:
25         classifier, cm = cm[0], cm[1]
26         cal_evaluation(classifier, cm)

```

```

1 # Confusion matrix, accuracy, precision and recall for random forest and logistic regression
2 confusion_matrices = [
3     ("Random Forest", confusion_matrix(y_test, best_RF_model.predict(X_test))),
4     ("Logistic Regression", confusion_matrix(y_test, best_LR_model.predict(X_test))),
5     ("K nearest neighbor", confusion_matrix(y_test, best_KNN_model.predict(X_test)))
6 ]
7
8 draw_confusion_matrices(confusion_matrices)

```

Random Forest  
 Accuracy is: 0.8604  
 precision is: 0.8007518796992481  
 recall is: 0.41846758349705304

Logistic Regression  
 Accuracy is: 0.8084  
 precision is: 0.6056338028169014  
 recall is: 0.16895874263261296

K nearest neighbor  
 Accuracy is: 0.776  
 precision is: 0.2  
 recall is: 0.03339882121807466

## ▼ Part 3.4: Model Evaluation - ROC & AUC

RandomForestClassifier, KNeighborsClassifier and LogisticRegression have predict\_proba() function

### ▼ Part 3.4.1: ROC of RF Model

```

1 from sklearn.metrics import roc_curve
2 from sklearn import metrics
3
4 # Use predict_proba to get the probability results of Random Forest
5 y_pred_rf = best_RF_model.predict_proba(X_test)[: , 1]
6 fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)

```

```


1 best_RF_model.predict_proba(X_test)

```



```
array([[0.69937921, 0.30062079],
       [0.95239486, 0.04760514],
       [0.74105858, 0.25894142],
       ...,
       [0.85967291, 0.14032709],
       [0.95083241, 0.04916759],
       [0.90647722, 0.09352278]])
```

```
1 # ROC curve of Random Forest result
2 import matplotlib.pyplot as plt
3 plt.figure(1)
4 plt.plot([0, 1], [0, 1], 'k--')
5 plt.plot(fpr_rf, tpr_rf, label='RF')
6 plt.xlabel('False positive rate')
7 plt.ylabel('True positive rate')
8 plt.title('ROC curve - RF model')
9 plt.legend(loc='best')
10 plt.show()
```



```
1 from sklearn import metrics
2
3 # AUC score
4 metrics.auc(fpr_rf, tpr_rf)
```

```
0.8457044914295074
```


### ▼ Part 3.4.1: ROC of LR Model

```
1 # Use predict_proba to get the probability results of Logistic Regression
2 y_pred_lr = best_LR_model.predict_proba(X_test)[: , 1]
3 fpr_lr, tpr_lr, thresh = roc_curve(y_test, y_pred_lr)
```

```
1 best_LR_model.predict_proba(X_test)
```

```
array([[0.83732252, 0.16267748],
       [0.91667203, 0.08332797],
       [0.85713621, 0.14286379],
       ...,
       [0.72754602, 0.27245398],
       [0.88726134, 0.11273866],
       [0.86330938, 0.13669062]])
```

```
1 # ROC Curve
2 plt.figure(1)
3 plt.plot([0, 1], [0, 1], 'k--')
4 plt.plot(fpr_lr, tpr_lr, label='LR')
5 plt.xlabel('False positive rate')
6 plt.ylabel('True positive rate')
7 plt.title('ROC curve - LR Model')
8 plt.legend(loc='best')
9 plt.show()
```



```
1 # AUC score
2 metrics.auc(fpr_lr, tpr_lr)
```

```
0.7703625055381831
```

## ▼ Part 4: Model Extra Functionality

### ▼ Part 4.1: Logistic Regression Model

The correlated features that we are interested in

```

1 X_with_corr = X.copy()
2
3 X_with_corr = OneHotEncoding(X_with_corr, enc_ohe, ['Geography'])
4 X_with_corr['Gender'] = enc_oe.transform(X_with_corr[['Gender']])
5 X_with_corr['SalaryInRMB'] = X_with_corr['EstimatedSalary'] * 6.4
6 X_with_corr.head()

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: FutureWarning:   
 warnings.warn(msg, category=FutureWarning)

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	0.0	42	2	0.00	1	1	1
1	608	0.0	41	1	83807.86	1	0	1
2	502	0.0	42	8	159660.80	3	1	0
3	699	0.0	39	1	0.00	2	0	0
4	850	0.0	43	2	125510.82	1	1	1



```

1 # add L1 regularization to logistic regression
2 # check the coef for feature selection
3 scaler = StandardScaler()
4 X_l1 = scaler.fit_transform(X_with_corr)
5 LRmodel_l1 = LogisticRegression(penalty="l1", C = 0.04, solver='liblinear')
6 LRmodel_l1.fit(X_l1, y)
7
8 indices = np.argsort(abs(LRmodel_l1.coef_[0]))[::-1]
9
10 print ("Logistic Regression (L1) Coefficients")
11 for ind in range(X_with_corr.shape[1]):
12     print ("{0} : {1}".format(X_with_corr.columns[indices[ind]], round(LRmodel_l1.coef_[0][ind], 4)))

```

Logistic Regression (L1) Coefficients

Age : 0.7307

IsActiveMember : -0.5046

Geography\_Germany : 0.3121

Gender : -0.2409

Balance : 0.151

CreditScore : -0.0457

NumOfProducts : -0.0438

Tenure : -0.0271

EstimatedSalary : 0.0055

Geography\_France : -0.0043

SalaryInRMB : 0.0037

HasCrCard : -0.0022

Geography\_Spain : 0.0

```

1 # add L2 regularization to logistic regression
2 # check the coef for feature selection
3 np.random.seed()
4 scaler = StandardScaler()
5 X_l2 = scaler.fit_transform(X_with_corr)
6 LRmodel_l2 = LogisticRegression(penalty="l2", C = 0.1, solver='liblinear', random_state=42)
7 LRmodel_l2.fit(X_l2, y)
8 LRmodel_l2.coef_[0]
9
10 indices = np.argsort(abs(LRmodel_l2.coef_[0]))[::-1]
11
12 print ("Logistic Regression (L2) Coefficients")
13 for ind in range(X_with_corr.shape[1]):
14     print ("{0} : {1}".format(X_with_corr.columns[indices[ind]],round(LRmodel_l2.coef_[0][in

```

Logistic Regression (L2) Coefficients

Age : 0.751

IsActiveMember : -0.5272

Gender : -0.2591

Geography\_Germany : 0.2279

Balance : 0.162

Geography\_France : -0.1207

Geography\_Spain : -0.089

CreditScore : -0.0637

NumOfProducts : -0.0586

Tenure : -0.0452

HasCrCard : -0.0199

SalaryInRMB : 0.0137

EstimatedSalary : 0.0137

## ▼ Part 4.2: Random Forest Model - Feature Importance Discussion

```

1 X_RF = X.copy()
2
3 X_RF = OneHotEncoding(X_RF, enc_ohe, ['Geography'])
4 X_RF['Gender'] = enc_oe.transform(X_RF[['Gender']])
5
6 X_RF.head()

```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: F  
warnings.warn(msg, category=FutureWarning)
```

```
CreditScore  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember
```

```
1 # check feature importance of random forest for feature selection
2 forest = RandomForestClassifier()
3 forest.fit(X_RF, y)
4
5 importances = forest.feature_importances_
6
7 indices = np.argsort(importances)[::-1]
8
9 # Print the feature ranking
10 print("Feature importance ranking by Random Forest Model:")
11 for ind in range(X.shape[1]):
12     print ("{0} : {1}".format(X_RF.columns[indices[ind]],round(importances[indices[ind]], 4)
```

Feature importance ranking by Random Forest Model:

Age : 0.2372

EstimatedSalary : 0.1483

Balance : 0.1429

CreditScore : 0.1426

NumOfProducts : 0.1302

Tenure : 0.0822

IsActiveMember : 0.0396

Geography\_Germany : 0.0217

HasCrCard : 0.0186

Gender : 0.0182