

✓ Credit Card Fraud Detection Model

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns, numpy as np
3 from sklearn.ensemble import IsolationForest
4 from numpy import genfromtxt
5 from scipy.stats import multivariate_normal
6 from sklearn.metrics import f1_score
7 import io
8 import pandas as pd
9 from sklearn.model_selection import train_test_split
10 from sklearn.metrics import f1_score, roc_auc_score, roc_curve, precision_recall_curve, auc, make_scorer, recall_score, accuracy_score, pr
11
```

```
1 creditcardDF = pd.read_csv('/content/drive/My Drive/fraudData/creditcard.csv')
2 creditcardDF.head()#all numerical
```

```
➡
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V1
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.99139
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.48909
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	0.207643	0.624501	0.066084	0.71729
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.50775
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.34585

```
1 #Distribution of the label column
2 creditcardDF['Class'].value_counts()
```

```
➡ 0    284315
   1     492
   Name: Class, dtype: int64
```

```
1 492 / (284315 + 492) #<0.2 percent
```

```
➡ 0.001727485630620034
```

```
1 creditcardDF.shape
```

```
➡ (284807, 31)
```


```
1 creditcardDF.isna().sum()#null checking
```

```
➡ Time      0
   V1      0
   V2      0
   V3      0
   V4      0
   V5      0
   V6      0
   V7      0
   V8      0
   V9      0
  V10      0
  V11      0
  V12      0
  V13      0
  V14      0
  V15      0
  V16      0
  V17      0
  V18      0
  V19      0
 V20      0
 V21      0
 V22      0
 V23      0
 V24      0
 V25      0
 V26      0
```

```
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```


```
1
2
3 sns.distplot(creditcardDF['Time'])
```

↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f4f248509b0>



```
1
2 sns.distplot(creditcardDF['Amount'])
```

↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f4f24066a90>



```
1 creditcardDF['Amount'] = np.log(creditcardDF['Amount'] + 1)
2
3 creditcardDF['Time'] = np.log(creditcardDF['Time'] + 1)
4
5 normal = creditcardDF[creditcardDF['Class'] == 0]
6 anomaly = creditcardDF[creditcardDF['Class'] == 1]
7
8 train, small_normal = train_test_split(normal, test_size=0.2, random_state=0)
9 normal_valid, normal_test = train_test_split(small_normal, test_size=0.5, random_state=0)#good hd
10 anomaly_valid, anomaly_test = train_test_split(anomaly, test_size=0.5, random_state=0)#10 bad hd
11
12 validation = pd.concat([normal_valid, anomaly_valid])#include both good and bad, cross validation data in our text
13 test = pd.concat([normal_test, anomaly_test])#include both good and bad
14
15
16 print(validation.shape)
17 print(test.shape)
18
19 #reset_index
20 train = train.drop(columns = ['Class']).reset_index(drop= True)#no need of label in train data, drop it
21 print(train.shape)
```

↗ (28677, 31)
(28678, 31)
(227452, 30)

```
1 featureNames = list(train.columns.values)#feature names only, no label
2 valFeatures = validation[featureNames].reset_index(drop= True)#feature df only, no label, will delete the current (sampled) index instead
3 testFeatures = test[featureNames].reset_index(drop= True)#feature df only, no label
4
5 valLabel = validation['Class']#label df only
```

```
6 testLabel = test['Class']#label df only
7
```

```
1 valFeatures.head()
```

<

```
1 validation['Class'].value_counts()
```


```
0    28431
1      246
Name: Class, dtype: int64
```

```
1 test['Class'].value_counts()
```


```
2
3
4
```

```
0    28432
1      246
Name: Class, dtype: int64
```


```
1 plt.figure()
2 plt.title("train data of V1 and V2")#may contain outliers
3 plt.xlabel("V1")
4 plt.ylabel("V2")
5 plt.plot(train.iloc[:, 1],train.iloc[:,2],"bx")
6 plt.show()
7
8
9
```



```
1 plt.figure()
2 plt.title("validation data of V1 and V2")
3 plt.xlabel("V1")
4 plt.ylabel("V2")
5 plt.plot(validation.iloc[:, 1],validation.iloc[:,2],"bx")
6 plt.show()
```



```
1 plt.figure()
2 plt.title("test data of V1 and V2")
3 plt.xlabel("V1")
4 plt.ylabel("V2")
5 plt.plot(test.iloc[:, 1], test.iloc[:, 2], "bx")
6 plt.show()
```



```
1 # np.arange(1, 20, 2)
```

```
array([ 1,  3,  5,  7,  9, 11, 13, 15, 17, 19])
```

```
1 #find parameter for each col/feature in df for the Gaussian distribution
2 def estimateGaussian(dataset):
3     mu = np.mean(dataset, axis=0)#vector
4     sigma = np.cov(dataset.T)#matrix
5     return mu, sigma
6
```

```
1
2 pdfVal = model.pdf(valFeatures)
3 print(max(pdfVal))#too small, can not differentiate
4 print(min(pdfVal))
5
6 p_val = model.logpdf(valFeatures)#Log of the pdf first, then apply to features, to change the magnitude of prob
7 print(max(p_val))#
8 print(min(p_val))
```

```
3.936022689247968e-12
0.0
-26.26085037221045
-7554.270217704667
```

```
1 p = model.logpdf(train)
2 print(p.shape)
3 print((p_val.shape))
4
5 # print(p)
6
```

```
(227452,)
(28677,)
```

```
1 print(p_val)
2 print(p_val < -500)
```

```
→ [ -31.28574735 -34.94205051 -27.79402451 ... -5175.93656039
    -4545.5057626 -29.7152192 ]
   [False False ... True True False]
```

```
1 [[1],[2],[3]]
2
3 ravel ->
4 [1,2,3]
```

```
1
2 #get the score list
3 scores = []
4 p_val = model.logpdf(valFeatures)#Log of the pdf
5
6 # thresholds = np.linspace(-1000, -10,150)
7 thresholds = np.linspace(min(p_val), max(p_val),200)#generate all candidate threshold, epsilon
8
9 #find optimal threshold: bestThreshold
10 for threshold in thresholds:
11     y_pred = (p_val < threshold).astype(int)# list of 0 and 1
12     #calculate recall, precision and f1 for each (truth, pred) pair, corresponding to that threshold
13     scores.append([recall_score(valLabel, y_pred),
14                   precision_score(valLabel, y_pred),
15                   f1_score(valLabel, y_pred, average = "binary")])
16
17 scores = np.array(scores)
18 maxIndex = scores[:,2].ravel().argmax()#maxIndex of the 3rd column (f1_score) #193, #.ravel return a flattened array
19 bestThreshold = thresholds[maxIndex]
20 print(scores.shape)#each row is a pair of (recall, precision, f1) corresponding to a threshold
```

```
→ /usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision is ill-defined and bei
   _warn_prf(average, modifier, msg_start, len(result))
   (200, 3)
```

```
1 print(scores)
```

```
→ [0.00813008 1.      0.01612903]
   [0.01219512 1.      0.02409639]
   [0.01219512 1.      0.02409639]
   [0.01219512 1.      0.02409639]
   [0.01219512 1.      0.02409639]
```

```

[0.1300813 1. 0.23021583]
[0.1300813 1. 0.23021583]
[0.1300813 1. 0.23021583]
[0.1300813 1. 0.23021583]
[0.1300813 1. 0.23021583]
[0.13414634 1. 0.23655914]
[0.15853659 1. 0.27368421]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]
[0.17479675 1. 0.29757785]

```

```

1 print(maxIndex)
2 print(bestThreshold)

```

```

→ 193
-253.2360071762041

```

```

1 np.mean(train.iloc[:,1])

```

```

→ 0.005246755420062154

```

```

1 mu[1]

```

```

→ 0.005246755420062154

```

```

1 print(mu)
2 # print(sigma)
3

```

```

→ Time      11.252384
V1          0.005247
V2         -0.005416
V3          0.010293
V4         -0.008144
V5          0.004281
V6          0.001813
V7          0.010354
V8         -0.001103
V9          0.006351
V10         0.009573
V11        -0.007736
V12         0.009943
V13         0.001084
V14         0.010816
V15         0.001082
V16         0.007216
V17         0.012364
V18         0.003412
V19        -0.001811
V20        -0.001092
V21        -0.001302
V22        -0.000354
V23         0.000209
V24         0.000288
V25         0.000375
V26         0.000457
V27        -0.000509
V28        -0.000119
Amount      3.152259
dtype: float64

```

```

1 #performance on test data
2 #prediction on test data
3 y_test_pred_raw = model.logpdf(testFeatures)
4 y_pred_test = y_test_pred_raw < bestThreshold
5
6 f1_score(testLabel, y_pred_test, average = "binary")

```

```

→ 0.7401574803149606

```

```
1 y_pred_test
```

```
→ array([False, False, False, ..., True, False, True])
```

```
1 #index of predicted outliers in test data
```

```
2 predoutliersTest = np.asarray(np.where(y_pred_test))
```

```
3
```

```
4 len(predoutliersTest[0])
```

```
→ 262
```

```
1 predoutliersTest #indexes
```

```
→ array([[ 248,   437,   605,  1007,  1353,  1451,  1462,  1546,  1988,
          2461,  3674,  3928,  4216,  4928,  5144,  5846,  5975,  6022,
          6682,  6706,  6858,  7017,  7138,  8267,  8452,  8611,  8677,
          8936,  8996,  9207,  9443,  9807,  9988, 10263, 10391, 10657,
         11224, 12205, 13539, 13935, 14050, 14573, 14579, 14802, 14869,
         15740, 16061, 16888, 17322, 17663, 19352, 19902, 20680, 20800,
         21748, 22366, 22552, 22859, 23217, 23456, 23742, 24639, 24819,
         25654, 25678, 26035, 27282, 27293, 27314, 27587, 27723, 28117,
         28178, 28396, 28432, 28433, 28434, 28435, 28436, 28437, 28438,
         28440, 28443, 28444, 28445, 28446, 28447, 28449, 28450, 28453,
         28454, 28455, 28456, 28457, 28458, 28459, 28460, 28461, 28462,
         28463, 28464, 28465, 28466, 28467, 28468, 28469, 28470, 28471,
         28472, 28473, 28475, 28479, 28480, 28481, 28482, 28483, 28484,
         28486, 28487, 28490, 28492, 28493, 28494, 28496, 28497, 28498,
         28499, 28500, 28501, 28502, 28503, 28505, 28506, 28507, 28508,
         28510, 28511, 28512, 28513, 28517, 28521, 28523, 28525, 28526,
         28527, 28528, 28529, 28530, 28531, 28532, 28536, 28538, 28539,
         28540, 28542, 28543, 28544, 28546, 28547, 28549, 28550, 28551,
         28552, 28553, 28554, 28555, 28556, 28558, 28559, 28560, 28561,
         28562, 28564, 28565, 28566, 28567, 28568, 28570, 28572, 28574,
         28575, 28576, 28577, 28578, 28579, 28580, 28581, 28583, 28584,
         28585, 28586, 28588, 28589, 28591, 28592, 28594, 28596, 28598,
         28599, 28600, 28601, 28602, 28603, 28604, 28605, 28606, 28607,
         28609, 28610, 28612, 28615, 28617, 28618, 28619, 28620, 28621,
         28622, 28623, 28625, 28626, 28628, 28629, 28630, 28631, 28632,
         28633, 28636, 28637, 28638, 28639, 28640, 28641, 28642, 28643,
         28645, 28646, 28647, 28648, 28649, 28650, 28651, 28652, 28653,
         28654, 28656, 28657, 28658, 28659, 28660, 28661, 28662, 28663,
         28664, 28666, 28669, 28670, 28671, 28672, 28673, 28674, 28675,
         28677]])
```

```
1 #outliers identified on test data feature column V2 V3
```

```
2 plt.figure()
```

```
3 plt.title("test_data with outlier flaged red")
```


```
4 plt.xlabel("V2")
```

```
5 plt.ylabel("V3")
```

```
6 plt.plot(testFeatures.iloc[:, 2],testFeatures.iloc[:,3],"bx")
```

```
7 plt.plot(testFeatures.iloc[predoutliersTest[0],1],testFeatures.iloc[predoutliersTest[0],2],"ro")
```

```
8 plt.show()
```



```
1 # generate evaluation metrics
```

```
2 print("%s: %r" % ("accuracy_score is: ", accuracy_score(testLabel, y_pred_test)))
```

```
3 print("%s: %r" % ("roc_auc_score is: ", roc_auc_score(testLabel, y_test_pred_raw)))#correction: should be y_pred_test instead of y_test_p
```

```
4 print("%s: %r" % ("f1_score is: ", f1_score(testLabel, y_pred_test )))#string to int
```

```
5
```

```
6 print ("confusion_matrix is: ")
```

```
7 cm = confusion_matrix(testLabel, y_pred_test)
```

```
8 cmDF = pd.DataFrame(cm, columns=['pred_0', 'pred_1'], index=['true_0', 'true_1'])
```

```
9 print(cmDF)
```

```

10 print('recall =',float(cm[1,1])/(cm[1,0]+cm[1,1]))
11 print('precision =', float(cm[1,1])/(cm[1,1] + cm[0,1]))#1.0
12
13

```

```

→ accuracy_score is: : 0.9953971685612665
roc_auc_score is: : 0.03898289914947546
f1_score is: : 0.7401574803149606
confusion_matrix is:
      pred_0  pred_1
true_0   28358     74
true_1     58    188
recall = 0.7642276422764228
precision = 0.7175572519083969

```

```

1 # convert 1/-1 to 0/1 for f1 calculation
2 def convert(x):
3     if x == 1:
4         return 0
5     else:
6         return 1
7 pred = IFModel.predict(testFeatures) #1 for inliers, -1 for outliers.
8 pred2 = list(map(convert, pred))
9 # pred2
10 import collections
11
12 counter=collections.Counter(pred2)
13 print(counter)#
14
15 f1_score(testLabel, pred2, average = "binary")#0.48 when added parameters

```

```

→ Counter({0: 28288, 1: 390})
0.48113207547169806

```